

## COSC 404 Database System Implementation

### Concurrency Control

**Dr. Ramon Lawrence**  
University of British Columbia Okanagan  
ramon.lawrence@ubc.ca

COSC 404 - Dr. Ramon Lawrence

## Concurrency Control Overview

**Concurrency control (CC)** is a mechanism for guaranteeing that concurrent transactions in the database exhibit the ACID properties. Specifically, the isolation property.

There are different concurrency control protocols:

- ◆ lock-based protocols, timestamp protocols, validation protocols

Page 2

COSC 404 - Dr. Ramon Lawrence

## Lock-Based Protocols

A **lock** is a mechanism to control concurrent access to data.

- ◆ An item can only be accessed through the lock.

Data items can be locked in two modes:

- ◆ **exclusive (X) mode**: Data item can be both read as well as written. X-lock is requested using **lock-X** instruction.
- ◆ **shared (S) mode**: Data item can only be read. S-lock is requested using **lock-S** instruction.

Lock requests are made to the concurrency control manager. A transaction can only proceed after the request is *granted* and must follow the restrictions of the lock.

Page 3

COSC 404 - Dr. Ramon Lawrence

## Lock-Based Protocols (2)

Lock-compatibility matrix:

	S	X
S	true	false
X	false	false

A transaction may be granted a lock on an item if the requested lock is *compatible* with locks already held on the item by other transactions.

- ◆ Any # of transactions can hold shared locks on an item.
- ◆ If any transaction holds an exclusive lock on the item, no other transaction may hold any lock on the item.
- ◆ If a lock cannot be granted, the requesting transaction is made to wait until all incompatible locks held by other transactions are released. The lock is then granted.

Page 4

COSC 404 - Dr. Ramon Lawrence

## Lock-Based Protocol Example

Example of a transaction performing locking:

```
lock-S(A);
read (A);
unlock(A);
lock-S(B);  ⇐ Another transaction updates B here.
read (B);
unlock(B);
display(A+B)
```

Simple locking is not sufficient to guarantee serializability.

- ◆ If A and B get updated in-between the read of A and B, the displayed sum would be wrong.
- ◆ A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks. Locking protocols restrict the set of possible schedules.

Page 5

COSC 404 - Dr. Ramon Lawrence

## Pitfalls of Lock-Based Protocols

Consider the partial schedule:

$T_3$	$T_4$
lock-X(B) read(B) B ← B-50 write(B)	
lock-X(A)	lock-S(A) read(A) lock-S(B)

- ◆ Neither  $T_3$  nor  $T_4$  can make progress as executing **lock-S(B)** causes  $T_4$  to wait for  $T_3$  to release its lock on B, while executing **lock-X(A)** causes  $T_3$  to wait for  $T_4$  to release its lock on A.
- ◆ Such a situation is called a **deadlock**. To handle a deadlock one of  $T_3$  or  $T_4$  must be rolled back and its locks released.

Page 6

## Pitfalls of Lock-Based Protocols (2)

The potential for deadlock exists in most locking protocols. Deadlocks are a necessary evil.

**Starvation** is also possible if the concurrency control manager is badly designed. For example:

- ◆ A transaction may be waiting for an exclusive lock on an item, while a sequence of other transactions request and are granted a shared lock on the same item.
- ◆ The same transaction is repeatedly rolled back due to deadlocks.

The concurrency control manager can be designed to prevent starvation.

- ◆ For example, do not grant a shared lock if the item is exclusively locked or a transaction is waiting for a lock-X.

Page 7



## The Two-Phase Locking Protocol

**Two-Phase Locking (2PL)** ensures conflict-serializable schedules by requiring all locks be acquired before first unlock.

### Phase 1: Growing Phase

- ◆ transaction may obtain locks
- ◆ transaction may not release locks

### Phase 2: Shrinking Phase

- ◆ transaction may release locks
- ◆ transaction may not obtain locks

The protocol ensures serializability. It can be proved that the transactions can be serialized in the order of their *lock points* (i.e. the point where a transaction acquired its final lock). Page 8

## The Two-Phase Locking Protocol (2)

2PL *does not* ensure freedom from deadlocks.

- ◆ Cascading roll-back is also possible under two-phase locking.

**Conservative 2PL** is deadlock free as all locks must be pre-declared and allocated at transaction start time.

**Strict 2PL** prevents cascading rollback as a transaction holds all its exclusive locks until it commits/aborts.

- ◆ Thus, uncommitted data is locked and cannot be accessed.

**Rigorous 2PL** is even stricter as *all* locks are held till commit/abort. (also cascade free)

- ◆ Transactions can be serialized in the order that they commit.

Most database systems use strict or rigorous 2PL.

Page 9

## Lock Conversions

Increased concurrency is possible by allowing lock conversions.

- ◆ **Upgrade** - convert shared lock to exclusive lock
- ◆ **Downgrade** - convert exclusive lock to shared lock

For two-phase locking with lock conversions:

- ◆ Upgrades and lock acquires are allowed in growing phase.
- ◆ Downgrades and lock releases are in the shrinking phase.

This protocol ensures serializability, but still relies on the programmer to insert the various locking instructions.

Page 10

## Automatic Acquisition of Locks

A simple automated algorithm can place lock requests for a transaction  $T_i$  issuing the standard read/write instructions:

- ◆ The operation  $\text{read}(D)$  is processed as:
  - ⇒ if  $T_i$  has a lock on  $D$  then  $\text{read}(D)$  otherwise
  - ⇒ request a **lock-S** on  $D$  (may be necessary to wait for a **lock-X**)
  - ⇒ when **lock-S** request is granted, then  $\text{read}(D)$
- ◆ The operation  $\text{write}(D)$  is processed as:
  - ⇒ if  $T_i$  has a **lock-X** on  $D$  then  $\text{write}(D)$  otherwise
  - ⇒ if  $T_i$  has a **lock-S** on  $D$  then upgrade lock on  $D$  to **lock-X**
    - may have to wait for upgrade
  - ⇒ otherwise request a new **lock-X**
  - ⇒ finally  $\text{write}(D)$  when receive upgrade or new lock
- ◆ All locks are released after commit or abort.

Page 11

## Example on Auto Lock Insertion

Abbreviations:

- ◆ A transaction  $T_i$  requesting a **lock-S** on  $D$  is given as:  $sl_i(D)$ .
- ◆ A transaction  $T_i$  requesting a **lock-X** on  $D$  is given as:  $xl_i(D)$ .
- ◆ A transaction  $T_i$  unlocking a data item  $D$  is given as:  $ul_i(D)$ .

Given the transaction, insert lock operations according to 2PL:

$T_1: r_1(A); r_1(C); w_1(B); w_1(C);$

Basic 2PL:

$sl_1(A); r_1(A); xl_1(C); r_1(C); xl_1(B); ul_1(A); w_1(B); ul_1(B); w_1(C); ul_1(C); c_1;$

locks may be released anytime after  
this operation when not needed

Page 12

## Example on Auto Lock Insertion (2)

Conservative 2PL:

$atomic(sl_1(A), xl_1(C), xl_1(B))$  ← locks may be released after they are no longer needed  
 $r_1(A); r_1(C); w_1(B); w_1(C); c_1; ul_1(A); ul_1(B); ul_1(C);$

Strict 2PL:

$sl_1(A); r_1(A); xl_1(C); r_1(C); xl_1(B); ul_1(A); w_1(B); w_1(C); c_1; ul_1(B); ul_1(C);$   
 ↑  
 read locks may be released before commit (after last lock operation)

Rigorous 2PL:

$sl_1(A); r_1(A); xl_1(C); r_1(C); xl_1(B); w_1(B); w_1(C); c_1; ul_1(A); ul_1(B); ul_1(C);$   
 ↗  
 all locks released after commit

Page 13

## Questions on 2PL

1) Given the following transactions, insert lock operations according to 2PL:

$T_1: r_1(A); w_1(A); r_1(B); w_1(B);$   
 $T_2: r_2(B); w_2(B); r_2(A); w_2(A);$

2) Write one non-serial schedule that obeys to 2PL, or argue why one is not possible.

3) Repeat #1 and #2 for these transactions:

$T_1: r_1(A); w_1(A); r_1(B); w_1(B); c_1$   
 $T_2: r_2(A); w_2(A); r_2(B); w_2(B); c_2$   
 $T_3: r_3(C); r_3(A); w_3(C); c_3$

Page 14

## Graph-Based Protocols

Graph-based protocols are an alternative to two-phase locking, but require additional rules on how transactions access the DB.


The simplest model is that we know the order in which database items are accessed:

- ◆ Impose a partial ordering  $\rightarrow$  on the set  $D = \{d_1, d_2, \dots, d_n\}$  of all data items.
- ◆ If  $d_i \rightarrow d_j$ , then any transaction accessing both  $d_i$  and  $d_j$  must access  $d_i$  before accessing  $d_j$ .
- ◆ Implies that the set  $D$  may now be viewed as a directed acyclic graph, called a *database graph*.
- ◆ Is ordering of data items realistic for a database system?

The **tree-protocol** is a simple kind of graph protocol.

Page 15

## Tree Protocol



Overview:

- ◆ Only exclusive locks are allowed.
- ◆ The first lock by  $T_i$  may be on any data item. Subsequently, a data  $Q$  can be locked by  $T_i$  only if the parent of  $Q$  is currently locked by  $T_i$ .
- ◆ Data items may be unlocked at any time.
- ◆ A data item cannot be relocked once it is unlocked.

Page 16

## Tree Protocol Discussion

The tree protocol ensures conflict serializability as well as freedom from deadlock.

Unlocking may occur earlier in the tree-locking protocol than in the two-phase locking protocol.

- ◆ shorter waiting times and increase in concurrency
- ◆ protocol is deadlock-free (no deadlock-related rollbacks)  
 ⇒ However, cascading rollbacks due to transaction aborts are possible.

However, a transaction may have to lock data items that it does not access.

- ◆ increased locking overhead and additional waiting time
- ◆ potential decrease in concurrency

Schedules not possible under two-phase locking are possible under tree protocol, and vice versa.

Page 17

## Timestamp-Based Protocol

Each transaction  $T_i$  is issued a timestamp  $TS(T_i)$  when it enters the system.

- ◆ If an **old** transaction  $T_i$  has timestamp  $TS(T_i)$ , a **new** transaction  $T_j$  has timestamp  $TS(T_j)$  where  $TS(T_i) < TS(T_j)$ .
- ◆ The timestamp can be assigned using the system clock or some logical counter that is incremented after every timestamp is assigned.

The protocol manages concurrent execution such that the timestamps determine the serializability order.

Timestamp protocols do not use locks, so deadlock cannot occur!

Page 18

## Timestamp-Based Protocol (2)

To ensure serializability, the protocol maintains for each data  $Q$  two timestamp values:

◆ **W-timestamp**( $Q$ ) is the largest timestamp of any transaction that executed **write**( $Q$ ) successfully.

◆ **R-timestamp**( $Q$ ) is the largest timestamp of any transaction that executed **read**( $Q$ ) successfully.

The timestamp ordering protocol ensures that any conflicting **read** and **write** operations are executed in timestamp order.

Page 19



## Timestamp-Based Protocols (3)

Suppose a transaction  $T_i$  issues a **read**( $Q$ ):

◆ If  $TS(T_i) < W\text{-timestamp}(Q)$ , then  $T_i$  needs to read a value of  $Q$  that was already overwritten.

⇒ Hence, the **read** operation is rejected, and  $T_i$  is rolled back.

◆ If  $TS(T_i) \geq W\text{-timestamp}(Q)$ , then the **read** operation is executed.

⇒ The  $R\text{-timestamp}(Q)$  is set to the maximum of  $R\text{-timestamp}(Q)$  and  $TS(T_i)$ .

Suppose that transaction  $T_i$  issues a **write**( $Q$ ):

◆ If  $TS(T_i) \geq R\text{-timestamp}(Q)$  AND  $TS(T_i) \geq W\text{-timestamp}(Q)$ , then the **write** operation is executed.

◆ If  $TS(T_i) < R\text{-timestamp}(Q)$ , then the value of  $Q$  that  $T_i$  is producing was previously read by newer transaction.

⇒ Hence, the **write** operation is rejected, and  $T_i$  is rolled back.

◆ If  $TS(T_i) < W\text{-timestamp}(Q)$ , then  $T_i$  is attempting to write an obsolete value of  $Q$ .  $T_i$  is rolled back.

Page 20

## Timestamp Example

A partial schedule for several data items for transactions with timestamps 1, 2, 3, 4, 5:

$T_1$	$T_2$	$T_3$	$T_4$	$T_5$
read(Y)	read(Y)			read(X)
		write(Y)		
	write(X) abort			read(Z)
read(X)		write(Z) abort		
				write(Y) write(Z)

Page 21

## Correctness of Timestamp-Ordering Protocol

The timestamp-ordering protocol guarantees serializability since all the arcs in the precedence graph are of the form:



Thus, there will be no cycles in the precedence graph.

Timestamp protocol ensures freedom from deadlock as no transaction ever waits.

But the schedule may not be cascade-free, and may not even be recoverable. (extensions are possible)

Page 22

## Thomas' Write Rule

Modified version of the timestamp-ordering protocol in which obsolete **write** operations may be ignored under certain circumstances:

◆ When  $T_i$  attempts to write data item  $Q$ , if  $TS(T_i) < W\text{-timestamp}(Q)$ , then  $T_i$  is attempting to write an obsolete value of  $\{Q\}$ . Hence, rather than rolling back  $T_i$  as the timestamp ordering protocol would have done, this **write** operation can be ignored. Otherwise protocol is unchanged.

**Thomas' Write Rule** allows greater potential concurrency. Unlike previous protocols, it allows some view-serializable schedules that are not conflict-serializable.

Page 23

## Questions on Timestamping

1) Indicate what happens during each of these schedules where concurrency control is performed using timestamps:

- $st_1; st_2; r_1(A); r_2(B); w_2(A); w_1(B);$
- $st_1; r_1(A); st_2; w_2(B); r_2(A); w_1(B);$
- $st_1; st_2; st_3; r_1(A); r_2(B); w_1(C); r_3(B); r_3(C); w_2(B); w_3(A);$
- $st_1; st_3; st_2; r_1(A); r_2(B); w_1(C); r_3(B); r_3(C); w_2(B); w_3(A);$

Page 24

## Multiple Granularity

To this point, we have been locking individual data items. It is beneficial to allow locking of various size data items.

- ◆ Define a hierarchy of data granularities, where the small granularities are nested within larger ones.
- ◆ Can be represented graphically as a tree.


When a transaction locks a node in the tree *explicitly*, it *implicitly* locks all the node's descendents in the same mode.

Granularity of locking (level in tree where locking is done):

- ◆ *fine granularity* (lower in tree): high concurrency, high locking overhead (e.g. record locking, attribute locking)
- ◆ *coarse granularity* (higher in tree): low locking overhead, low concurrency (e.g. table locking, database locking)

Page 25

## Example of Granularity Hierarchy



The highest level in the hierarchy is the entire database.

The levels below are *relation*, *tuple* and *field* in that order.

Page 26

## Intention Lock Modes

In addition to S and X lock modes, there are three additional lock modes with multiple granularity:

- ◆ **intention-shared (IS)**: indicates explicit locking at a lower level of the tree but only with shared locks.
- ◆ **intention-exclusive (IX)**: indicates explicit locking at a lower level with exclusive or shared locks
- ◆ **shared and intention-exclusive (SIX)**: the subtree rooted by that node is locked explicitly in shared mode and explicit locking is being done at a lower level with exclusive-mode locks.

Intention locks allow a higher level node to be locked in S or X mode without having to check all descendent nodes.

Page 27

## Compatibility Matrix with Intention Lock Modes

The compatibility matrix for all lock modes is:


	IS	IX	S	SIX	X
IS	✓	✓	✓	✓	×
IX	✓	✓	×	×	×
S	✓	×	✓	×	×
SIX	✓	×	×	×	×
X	×	×	×	×	×

Page 28

## Multi-granularity Locking

Strongest

Weakest



Page 29



## Multiple Granularity Locking

Transaction  $T_i$  can lock a node  $Q$  using the rules:


- ◆ The lock compatibility matrix must be observed.
- ◆ The root of the tree must be locked first (in any mode).
- ◆ A node  $Q$  can be locked by  $T_i$  in S or IS mode only if the parent of  $Q$  is currently locked by  $T_i$  in either IX or IS mode.
- ◆ A node  $Q$  can be locked by  $T_i$  in X, SIX, or IX mode only if the parent of  $Q$  is currently locked by  $T_i$  in either IX or SIX mode.
- ◆  $T_i$  can lock a node only if it has not previously unlocked any node (that is, this is a variant of two-phase locking).
- ◆  $T_i$  can unlock a node  $Q$  only if none of the children of  $Q$  are currently locked by  $T_i$ .

Observe that locks are acquired in root-to-leaf order, whereas they are released in leaf-to-root order.

Page 30

### Multiple Granularity Locking Example


T1 wants to lock R1.t2.f1 in X-mode.



Page 31

### Multiple Granularity Locking Example (2)

T2 wants to lock R1.t2.f2 in X-mode. Does it work?




Yes, it works!

Page 32

### Multiple Granularity Locking Example (3)

T2 wants to lock R1.t2.f2 in X-mode. Does it work?




No, conflict at t2!

Page 33

### Multiple Granularity Locking Example (4)

T2 wants to lock R1.t2.f2 in X-mode. Does it work?




Yes, it works!

Page 34

### Multiple Granularity Locking Example (5)

T2 wants to lock R1.t2.f2 in S-mode. Does it work?




Yes, it works!

Page 35

### Multiple Granularity Locking Example (6)

T2 wants to lock R1.t2.f2 in X-mode. Does it work?



No, conflict at R1!

Page 36

## Multiversion Schemes

Multiversion schemes keep old versions of data to increase concurrency. This is especially useful for read transactions.

- ◆ Implemented in Microsoft SQL Server (optimistic concurrency) and PostgreSQL.

Each successful **write** creates a new version of the data item. Use timestamps or transaction ids to label versions.

When a **read** operation is issued, select an appropriate version of the data item based on the timestamp.

**Reads** never have to wait as an appropriate version is returned immediately.

Page 37

## Multiversion Timestamp Ordering

Each data item  $Q$  has a sequence of versions  $\langle Q_1, Q_2, \dots, Q_m \rangle$ . Each version  $Q_k$  contains three data fields:

- ◆ **Content** - the value of version  $Q_k$
- ◆ **W-timestamp**( $Q_k$ ) - timestamp of the transaction that created (wrote) version  $Q_k$
- ◆ **R-timestamp**( $Q_k$ ) - largest timestamp of a transaction that successfully read version  $Q_k$

When a transaction  $T_i$  creates a new version  $Q_k$  of  $Q$ ,  $Q_k$ 's W-timestamp and R-timestamp are initialized to  $TS(T_i)$ .

R-timestamp of  $Q_k$  is updated whenever a transaction  $T_j$  reads  $Q_k$ , and  $TS(T_j) > R\text{-timestamp}(Q_k)$ .

Page 38

## Multiversion Timestamp Scheme

The following scheme ensures serializability:

- ◆ Let  $Q_k$  denote the version of  $Q$  whose write timestamp is the largest write timestamp less than or equal to  $TS(T_i)$ .

If transaction  $T_i$  issues a **read**( $Q$ ) then:

- ◆ The value returned is the content of version  $Q_k$ .

If transaction  $T_i$  issues a **write**( $Q$ ) operation:

- ◆ If  $TS(T_i) < R\text{-timestamp}(Q_k)$ , then  $T_i$  is rolled back.
- ◆ If  $TS(T_i) = W\text{-timestamp}(Q_k)$ ,  $Q_k$  is overwritten.
- ◆ Otherwise a new version of  $Q$  is created.

Page 39

## Multiversion Timestamp Scheme (2)

Reads always succeed; writes may be rejected if:

- ◆ Some other transaction  $T_j$  that (in the serialization order defined by the timestamp values) should read  $T_i$ 's write, has already read a version created by a transaction older than  $T_i$ .

Challenges:

- ◆ Must have an efficient way of handling versions (and discarding when no longer needed).
- ◆ Conflicts resolved through rollbacks rather than waiting so user application must be prepared to resubmit failed transactions.
  - ⇒ Only update transactions can be rolled back.

Page 40

## Insert and Delete Operations

In addition to read/write operations, the system must handle **delete** and **insert** operations.

Deletion with two-phase locking:

- ◆ May only be performed if the transaction deleting the tuple has an exclusive lock on the tuple to be deleted.

Insertion with two-phase locking:

- ◆ A transaction that inserts a new tuple into the database is given an X-mode lock on the tuple.

Page 41

## The Phantom Phenomenon

Inserts/deletes can lead to the **phantom phenomenon**:

- ◆ A transaction that scans a relation (e.g., find all students) and a transaction that inserts a tuple in the relation (e.g., inserts a new student) may conflict in spite of not accessing any tuple in common.
- ◆ If only tuple locks are used, non-serializable schedules can result: the scan transaction may not see the new tuple, yet may be serialized before the insert transaction.
- ◆ Transactions conflict over a **phantom tuple**.

The transaction scanning the relation reads information that indicates what tuples the relation contains. A transaction inserting a tuple updates the same info.

**This information should be locked.**

Page 42

## The Phantom Phenomenon (2)

Can prevent problem by:

- ◆ Accepting the issue (read committed isolation)
- ◆ Locking the entire relation (multi-granularity locking)
- ◆ Using index-locking or predicate-locking to guarantee that conflicts within the relation are detected.
- ◆ Having a special lock associated with the entire file. Read transactions that scan the whole relation must get a read lock on it and update transactions must get a write lock.

Page 43

## Transaction Definition in SQL

In SQL, a transaction begins implicitly.

A transaction in SQL ends by:

- ◆ **Commit** accepts updates of current transaction.
- ◆ **Rollback** aborts current transaction and discards its updates. Failures may also cause a transaction to be aborted.

An **isolation level** reflects how a transaction perceives the results of other transactions. It applies only to your perspective of the database, not other transactions/users. Lowering isolation level improves performance but may potentially sacrifice consistency.

Page 44

## Example Transactions

Transaction to deposit \$50 into a bank account:

```
BEGIN TRANSACTION;
UPDATE Account WHERE num = 'S1' SET balance=balance+50;
COMMIT T1;
```

Transaction to calculate totals for all accounts (twice):

```
BEGIN TRANSACTION;
SELECT SUM(balance) as total1 FROM Account;
SELECT SUM(balance) as total2 FROM Account;
COMMIT T2;
```

Transaction to add a new account:

```
BEGIN TRANSACTION;
INSERT INTO ACCOUNT (num, balance) VALUES ('S5', 100);
COMMIT T3;
```

Page 45



## Levels of Consistency in SQL-92

The isolation level can be specified by:

SET TRANSACTION ISOLATION LEVEL = X where X is

- ◆ **Serializable** - transactions behave like executed one at a time.
- ◆ **Repeatable read** - repeated reads must return same data. Does not necessarily read newly inserted records.
- ◆ **Read committed** - only committed values can be read, but successive reads may return different values.
- ◆ **Read uncommitted** - even uncommitted records may be read. Reading an uncommitted value is called a *dirty read*.

Page 46

## Scheduling of Transactions

Each transaction in a database is a separate executing program.

- ◆ A transaction may be its own program or a thread of execution.

The operating system schedules the execution of programs outside of the control of the DBMS.

- ◆ Thus, transactions may be executed in any order (as long as the order of operations within a transaction are the same). This interleaving is what produces different schedules.

The DBMS uses its concurrency control protocol to restrict the schedules to those that respect the consistency specified by the user for the transaction isolation level.

- ◆ All transactions must write lock any data item updated and the relation lock if inserting.
- ◆ Isolation level only affects read locks.

Page 47

## Isolation Example Serializable

A **serializable** schedule requires that regardless of the interleaving of the operations, the final result is the same as some serial ordering of the transactions.

- ◆ Thus, we do not have to test all possible schedules as it is sufficient to just enumerate the serial schedules.
- ◆ Read and write locks are held to commit. Also have a relation-level lock.

For three transactions, there are  $3! = 6$  serial schedules.

For these examples, assume that the total amount of money in all accounts is \$5000 before the transactions begin.

Page 48



## Isolation Example Serializable (2)

COSC 404 - Dr. Ramon Lawrence

Example schedule for T1, T2, T3:

```
UPDATE Account WHERE num = 'S1' SET balance=balance+50;
COMMIT T1;
SELECT SUM(balance) as total1 FROM Account;
SELECT SUM(balance) as total2 FROM Account;
COMMIT T2;
INSERT INTO ACCOUNT (num, balance) VALUES ('S5' , 100);
COMMIT T3;
```

After execution, total1 = \$5050 and total2 = \$5050.

◆The results for all six serial schedules are:

⇒ T1, T2, T3 – total1 = \$5050 ; total2 = \$5050  
 ⇒ T1, T3, T2 – total1 = \$5150 ; total2 = \$5150  
 ⇒ T2, T1, T3 – total1 = \$5000 ; total2 = \$5000  
 ⇒ T2, T3, T1 – total1 = \$5000 ; total2 = \$5000  
 ⇒ T3, T1, T2 – total1 = \$5150 ; total2 = \$5150  
 ⇒ T3, T2, T1 – total1 = \$5100 ; total2 = \$5100

Page 49

## Isolation Example Repeatable read

COSC 404 - Dr. Ramon Lawrence

With **repeatable read**, a transaction is guaranteed to get the same data back on multiple reads but may see **phantom records** inserted in between reads.

◆Read and write locks are held to commit.

Example schedule:

```
UPDATE Account WHERE num = 'S1' SET balance=balance+50;
COMMIT T1;
SELECT SUM(balance) as total1 FROM Account;
INSERT INTO ACCOUNT (num, balance) VALUES ('S5' , 100);
COMMIT T3;
SELECT SUM(balance) as total2 FROM Account;
COMMIT T2;
```

After execution, total1 = \$5050 and total2 = \$5150 as the second read sees the newly inserted tuple.

Page 50

## Isolation Example Read Committed

COSC 404 - Dr. Ramon Lawrence

With **read committed**, each read will get the most recently committed values even if different than an earlier read.

◆Read locks are released after every statement. Write locks released at commit.

Example schedule:

```
SELECT SUM(balance) as total1 FROM Account;
UPDATE Account WHERE num = 'S1' SET balance=balance+50;
COMMIT T1;
INSERT INTO ACCOUNT (num, balance) VALUES ('S5' , 100);
COMMIT T3;
SELECT SUM(balance) as total2 FROM Account;
COMMIT T2;
```

After execution, total1 = \$5000 and total2 = \$5150 as the second read sees the newly inserted tuple and T1's update.

Page 51

## Isolation Example Read Uncommitted

COSC 404 - Dr. Ramon Lawrence

Read uncommitted allows a transaction to read dirty data that has not been (and may never be) committed.

◆Transaction acquires no read locks.

Example schedule:

```
UPDATE Account WHERE num = 'S1' SET balance=balance+50;
SELECT SUM(balance) as total1 FROM Account;
INSERT INTO ACCOUNT (num, balance) VALUES ('S5' , 100);
SELECT SUM(balance) as total2 FROM Account;
COMMIT T2;
ABORT T3;
ABORT T1;
```

After execution, total1 = \$5050 and total2 = \$5150 as T2's sees even uncommitted data. Note that both T1 and T3 abort so T2 sees incorrect data. **It is very dangerous to use read uncommitted if the transaction updates the database!**

Page 52

## Summary of Isolation Levels

COSC 404 - Dr. Ramon Lawrence

Isolation Level	Problems	Lock Usage	Speed	Comments
Serializable	None	Read locks held to commit ; read lock on relation	Slowest	Only level that guarantees correctness.
Repeatable read	Phantom tuples	Read locks held to commit	Medium	Useful for modify transactions.
Read committed	Phantom tuples, values may change	Read locks released after each statement	Fast	Useful for transactions where operations are separable but updates are all or none.
Read uncommitted	Phantoms, values may change, dirty reads	No read locks	Fastest	Useful for read-only transactions that tolerate inaccurate results

Page 53

## Transaction Practice Question

COSC 404 - Dr. Ramon Lawrence

Given these transactions and table Bid(itemID, price) that initially contains the two tuples: (i1,10) and (i2,20):

```
T1: BEGIN TRANSACTION;
    S1: UPDATE Bid SET price = price + 5;
    S2: INSERT INTO Bid VALUES (i3,30);
    COMMIT;
```

```
T2: BEGIN TRANSACTION;
    S1: SELECT SUM(price) AS p1 FROM Bid;
    S2: SELECT MAX(price) AS p2 FROM Bid;
    COMMIT;
```

Assume that T1 executes with isolation level serializable and both transactions successfully commit.

◆1) If T2 executes with isolation level serializable, what are all the possible pairs of values for p1 and p2 returned by T2?

◆2) If T2 executes with isolation level read committed, what are all the possible pairs of values for p1 and p2 for T2?

Page 54

## Deadlock Handling

A system is deadlocked if there is a set of transactions such that every transaction in the set is waiting for another transaction in the set.

Two mechanisms for deadlock handling:

- ◆ **deadlock prevention** - do not allow system to enter deadlock state
- ◆ **deadlock detection** - detect deadlock condition and abort transactions to remove deadlock state

Cost of deadlock handling includes:

- ◆ overhead of scheme itself
- ◆ potential losses in transaction processing due to rollbacks

Page 55

## Deadlock Prevention

**Deadlock prevention** protocols ensure that the system will *never* enter into a deadlock state.

Some strategies:

- ◆ Require that each transaction locks all its data items before it begins execution (predeclaration).
- ◆ Impose partial ordering of all data items and require that a transaction can lock data items only in the order specified by the partial order (graph-based (tree) protocol).
- ◆ Others include wound-wait and wait-die strategies that use timestamps to determine transaction age and determine if a transaction should wait or be rolled back on a lock conflict.

Page 56

## Wound-Wait and Wait-Die Strategies

**Wait-Die** scheme — non-preemptive

- ◆ Older transaction may wait for younger one to release data item. Younger transactions never wait for older ones; they are rolled back instead.
- ◆ A transaction may die several times before acquiring needed data item.

**Wound-Wait** scheme — preemptive

- ◆ Older transaction *wounds* (forces rollback) of younger transaction instead of waiting for it. Younger transactions may wait for older ones.
- ◆ May cause fewer rollbacks than *wait-die* scheme.

Note: A rolled back transaction is restarted with its original timestamp. Older transactions have precedence over newer ones, and starvation is avoided.

Page 57

## Timeout-Based Schemes

In a Timeout-Based Schemes:

- ◆ A transaction waits for a lock only for a specified amount of time. After that, the transaction times out and is rolled back.
- ◆ Thus deadlocks are not possible.
- ◆ Simple to implement, but starvation is possible.
- ◆ Difficult to determine good value of the timeout interval.
  - ⇒ Too short - false deadlocks (unnecessary rollbacks)
  - ⇒ Too long - wasted time while system is in deadlock

Page 58

## Deadlock Detection & Recovery

If deadlocks are not prevented, then a detection and recovery procedure is needed to recover when the system enters the deadlock state.

An algorithm is run periodically to check for deadlock. If the system is in deadlock, then transactions are aborted to resolve the deadlock.

Deadlock detection requires the system:

- ◆ Maintain information about currently allocated locks.
- ◆ Provide an algorithm to detect a deadlock state.
- ◆ Recover from deadlock by aborting transactions efficiently.

Page 59



## Wait-for Graphs

Deadlocks can be detected using a **wait-for graph**, which consists of a pair  $G = (V, E)$ :

- ◆  $V$  is a set of vertices (all the transactions in the system).
- ◆  $E$  is a set of edges; each element is an ordered pair  $T_i \rightarrow T_j$ .
- ◆ If  $T_i \rightarrow T_j$  is in  $E$ , then there is a directed edge from  $T_i$  to  $T_j$ , implying that  $T_i$  is waiting for  $T_j$  to release a data item.


When  $T_i$  requests a data item currently being held by  $T_j$ , then the edge  $T_i \rightarrow T_j$  is inserted into the graph.

- ◆ This edge is removed only when  $T_j$  is no longer holding a data item needed by  $T_i$ .


The system is in a deadlock state if and only if the wait-for graph has a **cycle**. Must invoke a deadlock-detection algorithm periodically to look for cycles.

Page 60

## Wait-for Graph Examples



Wait-for graph with no cycle



Wait-for graph with a cycle

## Deadlock Recovery

When a deadlock is detected three factors to consider:

- ◆ **Victim selection** - Some transaction will have to rolled back (made a victim) to break deadlock.
  - ⇒ Select the victim transaction that will incur minimum cost (computation time, data items used, etc.).
- ◆ **Rollback** - determine how far to roll back transaction
  - ⇒ **Total rollback**: Abort the transaction and then restart it.
  - ⇒ More effective to roll back transaction only as far as necessary to break deadlock. (requires system store additional information)
- ◆ **Starvation** happens if same transaction is always chosen as victim.
  - ⇒ Include the number of rollbacks in the cost factor to avoid starvation.

## Questions on Deadlocks

1) Assume a read-lock is requested before each read, and a write lock before each write. All unlocks occur after the last operation of a transaction. Explain what operations are denied during each schedule, draw the wait-for graph, and pick a transaction to abort if a deadlock does occur.

- a)  $r_1(A)$ ;  $r_2(B)$ ;  $w_1(C)$ ;  $r_3(D)$ ;  $r_4(E)$ ;  $w_3(B)$ ;  $w_2(C)$ ;  $w_4(A)$ ;  $w_1(D)$ ;
- b)  $r_1(A)$ ;  $r_2(B)$ ;  $r_3(C)$ ;  $w_1(B)$ ;  $w_2(C)$ ;  $w_3(D)$ ;
- c)  $r_1(A)$ ;  $r_2(B)$ ;  $r_3(C)$ ;  $w_1(B)$ ;  $w_2(C)$ ;  $w_3(A)$ ;

## Concurrency Control Summary

**Concurrency control protocols** are used to ensure concurrent transactions maintain their isolation.

**Two-phase locking (2PL)** is a common protocol.

- ◆ Other protocols: timestamping, resource ordering, optimistic
- ◆ Multigranularity locking schemes are also possible.

**Multiversion schemes** create new versions on every update and determine the correct version for reads. They are often used for optimistic concurrency control.

**Deadlocks** must be handled by either deadlock prevention or deadlock detection and recovery.

- ◆ Prevention: wound-wait and wait-die schemes
- ◆ Detection: wait-for graphs and transaction rollback

## Major Objectives

The "One Things":

- ◆ Explain how two-phase locking (2PL) works and detect valid 2PL schedules.
- ◆ Perform deadlock detection and recovery using wait-for graphs.
- ◆ Explain and use the timestamp based protocol.
- ◆ Perform multiple granularity locking using lock modes, rules, and compatibility matrix.

Major Theme:

- ◆ Concurrency control allows only serializable schedules. Although many techniques have been developed, variations of 2PL are used in commercial databases. Locking methods must have a method of handling or preventing deadlocks.

## Objectives

- ◆ Define concurrency control, locking protocol, deadlock, starvation, exclusive and shared locks (compatibility matrix).
- ◆ Define and use conservative, strict, and rigorous 2PL.
- ◆ Explain the use of lock conversions (upgrades/downgrades).
- ◆ Insert locks into a schedule using automatic algorithm.
- ◆ Perform locking using tree protocol.
- ◆ Define and motivate a validation based protocol.
- ◆ Explain the motivation for multiversion 2PL and timestamping.
- ◆ List some methods for deadlock prevention.
- ◆ List three factors with deadlock recovery.
- ◆ Explain how the phantom phenomenon occurs.
- ◆ List consistency levels in SQL-92 and determine which schedules are valid under each consistency level.